

Start coding or [generate](#) with AI.

Double-click (or enter) to edit

```
# # Internship Task 1: Big Data Analysis using PySpark
# This notebook analyzes a large fraud detection dataset using PySpark to demonstrate scalability.
```

```
# 🛠 Step 1: Install Java (Spark needs Java runtime)
!apt-get install openjdk-11-jdk-headless -qq > /dev/null

# 🛠 Step 2: Download Spark 3.5.1 (you can change version if needed)
!wget -q https://archive.apache.org/dist/spark/spark-3.5.1/spark-3.5.1-bin-hadoop3.tgz

# 🛠 Step 3: Extract Spark package
!tar -xzf spark-3.5.1-bin-hadoop3.tgz

# 🛠 Step 4: Install findspark to connect Python with Spark
!pip install -q findspark
```


```
import os
import findspark

# Set environment paths for Java and Spark
os.environ["JAVA_HOME"] = "/usr/lib/jvm/java-11-openjdk-amd64"
os.environ["SPARK_HOME"] = "/content/spark-3.5.1-bin-hadoop3"

# Initialize findspark
findspark.init()
```

Double-click (or enter) to edit

```
from google.colab import files
uploaded = files.upload()
```



Choose Files

fraud_detec...dataset.csv

- fraud_detection_bank_dataset.csv**(text/csv) - 5010367 bytes, last modified: 8/8/2021 - 100% done
 Saving fraud_detection_bank_dataset.csv to fraud_detection_bank_dataset.csv

```
from pyspark.sql import SparkSession

# Start or get the SparkSession
spark = SparkSession.builder \
    .appName("My Fraud Detection Analysis") \
    .getOrCreate()

df = spark.read.csv("fraud_detection_bank_dataset.csv", header=True, inferSchema=True)
df.printSchema()
df.show(5)
```



What can I help you build?



```
-- col_84: integer (nullable = true)
-- col_85: integer (nullable = true)
-- col_86: integer (nullable = true)
-- col_87: integer (nullable = true)
-- col_88: integer (nullable = true)
-- col_89: integer (nullable = true)
-- col_90: integer (nullable = true)
-- col_91: integer (nullable = true)
-- col_92: integer (nullable = true)
-- col_93: integer (nullable = true)
-- col_94: integer (nullable = true)
-- col_95: integer (nullable = true)
-- col_96: integer (nullable = true)
-- col_97: integer (nullable = true)
-- col_98: integer (nullable = true)
-- col_99: integer (nullable = true)
-- col_100: integer (nullable = true)
-- col_101: integer (nullable = true)
-- col_102: integer (nullable = true)
-- col_103: integer (nullable = true)
-- col_104: integer (nullable = true)
-- col_105: integer (nullable = true)
-- col_106: integer (nullable = true)
-- col_107: integer (nullable = true)
-- col_108: integer (nullable = true)
-- col_109: integer (nullable = true)
-- col_110: integer (nullable = true)
-- col_111: integer (nullable = true)
-- targets: integer (nullable = true)
```

```
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|_c0|col_0|col_1|col_2|col_3|col_4|col_5|col_6|col_7|col_8|col_9|col_10|col_11|col_12|col_13|col_14|col_15|col_16|col_17|col_18|col_19|
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
| 0| 9| 1354| 0| 18| 0| 1| 7| 9| 0| 0| 0| 0| 0| 0| 1| 0| 1| 0| 0| 0|
| 1| 0| 239| 0| 1| 0| 1| 0| 0| 0| 0| 0| 0| 0| 0| 0| 0| 1| 0| 0| 0|
| 2| 0| 260| 0| 4| 0| 3| 6| 0| 0| 0| 0| 0| 0| 1| 1| 0| 0| 0| 0|
| 3| 17| 682| 0| 1| 0| 0| 8| 17| 0| 0| 0| 0| 0| 0| 0| 0| 1| 0| 0| 0|
| 4| 1| 540| 0| 2| 0| 1| 7| 1| 0| 0| 0| 0| 0| 0| 1| 0| 1| 0| 0| 0|
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
```

only showing top 5 rows





Close

```
print("Row count:", df.count())
print("Column count:", len(df.columns))
```



Row count: 20468
Column count: 114

```
df.describe().show()
```



```
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|summary|_c0|col_0|col_1|col_2|col_3|col_4|
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
| count| 20468| 20468| 20468| 20468| 20468| 20468| |
| mean| 10233.5| 3.2262556185264804| 294.79304279851476| 0.42002149697088137| 2.329343365253078| 0.08359390267735001| 0.93985733828|
| stddev| 5908.746990691005| 20.56430794451052| 717.5419842156925| 7.367275061894912| 10.06851230174707| 0.8405365704812978| 4.2228960613|
| min| 0| 0| 0| 0| 0| 0| 0|
| max| 20467| 2301| 37808| 904| 772| 54|
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
```

```
from pyspark.sql.functions import sum
```

```
# Get columns where all values = 0
zero_cols = []
for col_name in df.columns:
    if df.select(sum(col_name)).first()[0] == 0:
        zero_cols.append(col_name)
```

```
# Drop them
df_clean = df.drop(*zero_cols)
print("Dropped columns:", zero_cols)
```

↻ Dropped columns: ['col_8', 'col_9', 'col_10', 'col_11', 'col_12', 'col_18', 'col_19', 'col_20', 'col_21', 'col_35', 'col_51', 'col_52',

df_clean.columns

↻ ['col_50',
'col_54',
'col_55',
'col_56',
'col_57',
'col_58',
'col_59',
'col_60',
'col_61',
'col_62',
'col_63',
'col_64',
'col_65',
'col_66',
'col_67',
'col_68',
'col_69',
'col_72',
'col_73',
'col_74',
'col_75',
'col_76',
'col_77',
'col_78',
'col_79',
'col_80',
'col_81',
'col_82',
'col_83',
'col_84',
'col_85',
'col_86',
'col_87',
'col_88',
'col_89',
'col_90',
'col_91',
'col_92',
'col_93',
'col_94',
'col_95',
'col_96',
'col_97',
'col_98',
'col_99',
'col_100',
'col_101',
'col_102',
'col_103',
'col_104',
'col_105',
'col_106',
'col_107',
'col_108',
'col_109',
'col_110',
'col_111',
'targets']

```
new_names = [  
    "amount", "oldbalanceOrig", "newbalanceOrig", "type", # first few known fields  
    # Add generic names for the rest  
] + [f"feature_{i}" for i in range(len(df_clean.columns) - 5)] + ["isFraud"] # assuming last column is the target  
  
df_clean = df_clean.toDF(*new_names)
```

df_clean.show(5)

↻

	amount	oldbalanceOrig	newbalanceOrig	type	feature_0	feature_1	feature_2	feature_3	feature_4	feature_5	feature_6	feature_7	feature_8	feature_9
	0	9	1354	0	18	0	1	7	9	0	1	0	1	
	1	0	239	0	1	0	1	0	0	0	0	0	1	

2	0	260	0	4	0	3	6	0	0	1	1	0
3	17	682	0	1	0	0	8	17	0	0	0	1
4	1	540	0	2	0	1	7	1	0	1	0	1

only showing top 5 rows

```
print("Rows:", df_clean.count())
print("Columns:", len(df_clean.columns))
```

Rows: 20468
Columns: 99

```
df_clean.printSchema()
```

```
-- feature_38: integer (nullable = true)
-- feature_39: integer (nullable = true)
-- feature_40: integer (nullable = true)
-- feature_41: integer (nullable = true)
-- feature_42: integer (nullable = true)
-- feature_43: integer (nullable = true)
-- feature_44: integer (nullable = true)
-- feature_45: integer (nullable = true)
-- feature_46: integer (nullable = true)
-- feature_47: integer (nullable = true)
-- feature_48: integer (nullable = true)
-- feature_49: integer (nullable = true)
-- feature_50: integer (nullable = true)
-- feature_51: double (nullable = true)
-- feature_52: integer (nullable = true)
-- feature_53: integer (nullable = true)
-- feature_54: integer (nullable = true)
-- feature_55: integer (nullable = true)
-- feature_56: integer (nullable = true)
-- feature_57: integer (nullable = true)
-- feature_58: integer (nullable = true)
-- feature_59: integer (nullable = true)
-- feature_60: integer (nullable = true)
-- feature_61: integer (nullable = true)
-- feature_62: integer (nullable = true)
-- feature_63: integer (nullable = true)
-- feature_64: integer (nullable = true)
-- feature_65: integer (nullable = true)
-- feature_66: integer (nullable = true)
-- feature_67: integer (nullable = true)
-- feature_68: integer (nullable = true)
-- feature_69: integer (nullable = true)
-- feature_70: integer (nullable = true)
-- feature_71: integer (nullable = true)
-- feature_72: integer (nullable = true)
-- feature_73: integer (nullable = true)
-- feature_74: integer (nullable = true)
-- feature_75: integer (nullable = true)
-- feature_76: integer (nullable = true)
-- feature_77: integer (nullable = true)
-- feature_78: integer (nullable = true)
-- feature_79: integer (nullable = true)
-- feature_80: integer (nullable = true)
-- feature_81: integer (nullable = true)
-- feature_82: integer (nullable = true)
-- feature_83: integer (nullable = true)
-- feature_84: integer (nullable = true)
-- feature_85: integer (nullable = true)
-- feature_86: integer (nullable = true)
-- feature_87: integer (nullable = true)
-- feature_88: integer (nullable = true)
-- feature_89: integer (nullable = true)
-- feature_90: integer (nullable = true)
-- feature_91: integer (nullable = true)
-- feature_92: integer (nullable = true)
-- feature_93: integer (nullable = true)
-- isFraud: integer (nullable = true)
```

```
df_clean.groupBy("isFraud").count().show() #count farud transaction
```

```
+-----+-----+
|isFraud|count|
+-----+-----+
|      1| 5438|
```

```
|      0|15030|
+-----+-----+
```

```
df_clean.filter("isFraud == 1").groupBy().avg("amount").show() # avg transation amount for fraud
```

```
↩ +-----+
  |avg(amount)|
  +-----+
  |      2718.5|
  +-----+
```

```
df_clean.groupBy("type", "isFraud").count().orderBy("type").show() # fraud count by transaction
```

```
↩ +---+-----+-----+
  |type|isFraud|count|
  +---+-----+-----+
  |  0|      0|12826|
  |  0|      1| 5276|
  |  1|      0| 1119|
  |  1|      1|   102|
  |  2|      1|    25|
  |  2|      0|   611|
  |  3|      1|     8|
  |  3|      0|   170|
  |  4|      0|    66|
  |  4|      1|    15|
  |  5|      1|     4|
  |  5|      0|   40|
  |  6|      0|   40|
  |  6|      1|     2|
  |  7|      1|     1|
  |  7|      0|   26|
  |  8|      0|   15|
  |  9|      0|   17|
  | 10|      0|    9|
  | 11|      0|    8|
  +---+-----+-----+
  only showing top 20 rows
```

Double-click (or enter) to edit

Insights:

- Transaction Type 0 had the highest volume and most frauds.

- Fraud was present across nearly all types.

- Type 1 had a significant fraud rate.

- 15 columns were dropped due to containing only 0s.

This shows the dataset had redundant information, and PySpark allowed fast filtering and analysis even for large structured data.